

Semantic Prioritized Experience Replay: VLM-Guided Posterior Reweighting for Sparse-Reward Robotic Manipulation

Daniel Grant

UC Berkeley, CS285 Deep Reinforcement Learning, Spring 2026

Abstract

We introduce **Semantic Prioritized Experience Replay (Semantic PER)**, a credit-assignment scheme in which a frozen vision-language model (VLM) serves as an exogenous posterior over which transitions in a failed episode are causally responsible for the failure. Concretely, the VLM emits an approximate posterior $q_\phi(t^*|\tau)$ over outcome-causing timesteps; we multiply the standard PER priority by a window kernel centred on this posterior, yielding the sampling distribution $\mu_{\text{Sem}}(i) \propto \mu_p(i) \cdot w_{\text{Sem}}(i;\tau)$. This factorisation locates the VLM as a credit-assignment oracle in an importance-sampling taxonomy alongside Retrace, V-trace and prioritised DQN. Empirically, on the Gymnasium-Robotics Fetch suite (Push, PickAndPlace, Slide) Semantic PER recovers a non-trivial fraction of the privileged-Oracle ceiling over PER+HER. A 12-episode cross-task pilot shows that a single prompt template produces 12/12 task-relevant failure annotations across three structurally different Fetch envs, evidence that VLM-derived priority signals are *task-agnostic by construction* in a way learned heads are not. We further propose **VLM-Verified Counterfactual Hindsight**, in which a counterfactual action proposed by the VLM is admitted to the buffer only after the simulator's own sparse reward fires on a fresh rollout from the failure state, eliminating an observed teleport-collapse failure mode by construction. Code, configs, and figures are at the project repository.

1. Introduction

Off-policy reinforcement learning on sparse-reward manipulation tasks is bottlenecked by *credit assignment*: a single 50-step trajectory ends with a binary success/failure signal, and the learner must work out which subset of decisions caused the outcome. Prioritised Experience Replay [1] addresses this with a TD-error proposal, $\mu_p(i) \propto |\delta_1|^{\alpha}$, and Hindsight Experience Replay (HER) [3] addresses it with achieved-future relabelling. Both schemes are *endogenous*: the credit signal is computed from the agent's own value function or trajectory. We ask whether a frozen foundation model can serve as an *exogenous* credit oracle that imports human-distilled physical priors over what typically causes manipulation failures.

The closest concurrent work is VLM-Guided Replay (Sharony et al. 2026, VLM-RB) [2], which prompts a VLM for goal-satisfaction on 32-frame sub-trajectory clips and linearly mixes the resulting score with uniform sampling. Our scheme differs along four orthogonal axes summarised in Section 2: (i) we localise the *cause of failure* rather than the presence of success, (ii) we operate on a per-timestep window rather than a 32-frame clip, (iii) we apply a *multiplicative* boost on top of PER rather than a linearly mixed component, and (iv) we measure VLM-attributable headroom against a privileged-state Oracle baseline.

Contributions. (1) We give the first principled framing of VLM-guided replay as *importance-sampled posterior reweighting* with a foundation-model credit-assignment oracle, locating our method in an off-policy correction taxonomy alongside Retrace, V-trace and prioritised DQN. (2) We propose **VLM-Verified Counterfactual Hindsight**, a simulator-grounded acceptance test that eliminates a previously observed teleport-collapse failure mode by construction. (3) We provide preliminary evidence of **cross-task transferability**: a single prompt template produces 12/12 task-relevant failure annotations across FetchPush, FetchPickAndPlace and FetchSlide with no per-task retraining.

2. Related Work

2.1 Goal-conditioned RL and HER

HER [3] is the canonical sparse-reward goal-conditioned method, with a growing literature of refinements: Next-Future [4] replaces HER's heuristic future strategy with a principled single-step rule; CEBP [5] prioritises by gripper-contact energy; HInt/NCII [6] uses counterfactual object-interaction nulling for relabelling; GCHR [7] adds hindsight regularisation; D-SPEAR [8] is a 2026 dual-stream PER for manipulation. Act2Goal [9] integrates HER with a learned world model. Our work targets the same sparse-reward setting but augments the priority surface rather than the relabelling distribution.

2.2 Prioritised replay and off-policy correction

PER [1] performs self-normalised importance sampling: a TD-error proposal μ_p with per-sample IS weight $w_{IS}(i) = (|D_B| \cdot \mu_p(i))^{-\beta}$. Retrace [10] and V-trace [11] generalise this to multi-step, per-step IS ratios with bounded variance; Freshness-Aware PER [12] addresses staleness in LM/VLM training. Section 3 locates our method inside this taxonomy.

2.3 VLMs in RL (and the VLM-RB comparison)

Foundation models enter RL in three flavours: as reward models (Rocamonde et al. [13], KAGI [14], VLM-CaR [15], Large Reward Models [16], MARVL, ERL-VLM); as critics/representations (Chen et al. [17], VLAC); and as hindsight/counterfactual relabellers (CAST [18], CoSo [19], AHA [20], AgentHER [21], ECHO [22]). VLM-RB [2] is the only prior published VLM-in-replay scheme; we differ in (i) failure-direction prompting, (ii) per-timestep localisation rather than 32-frame clip scoring, (iii) multiplicative TD-weight rather than additive mixture-with-uniform, and (iv) inclusion of a privileged-state Oracle headroom analysis. The differentiation matrix is summarised in Figure 1.

Method comparison: VLM-RB vs. our work

Axis	VLM-RB (Sharony et al., Feb 2026)	Ours (Path A + C)
Signal direction	Success / goal-satisfaction	Failure causation
Granularity	32-frame sub-trajectory clip score	Single causal timestep + bounded window
TD blending	Mixture-with-uniform; lambda warm-up 0 -> 0.5	Multiplicative weight on delta , no warm-up
VLM model	Frozen PerceptionLM (open-weights, async worker)	GPT-4o API + Claude (Path C); heuristic oracle
Benchmark	MiniGrid DoorKey + OGBench Scene 3/4/5 (UR5e)	Gymnasium-Robotics Fetch (Push, PickPlace, Slide)
Headroom analysis	None	Privileged-state Oracle quantifies VLM-attributable gap
Counterfactual reuse	Re-weights observed transitions only	Path C generates synthetic transitions from VLM goals

Figure 1. Differentiation against VLM-RB [2] along four orthogonal axes: signal direction (failure vs. success), granularity (per-timestep vs. 32-frame clip), prioritisation form (multiplicative vs. additive mixture), and the inclusion of an Oracle headroom baseline. See Section 5.4 for the cross-task transferability result.

3. Theoretical Motivation: VLM as Posterior

We argue that VLM-guided replay is best viewed as **importance-sampled stochastic optimisation with the VLM as a learned proposal distribution**. Specifically, the VLM emits a posterior $q_\phi(t^*|\tau)$ over which transition is responsible for an episode's outcome; semantic PER samples from a weighted product of TD-error priority and posterior mass.

3.1 Reweighting view of replay

Off-policy value learning solves a regression on a distribution μ over transitions, $L(\theta) = E_{i \sim \mu} [\mathbb{I}(\delta_i(\theta))]$. PER chooses $\mu = \mu_p$ with $\mu_p(i) \propto (|\delta_i| + \epsilon)^\alpha$, and the IS weight $w_{IS}(i) = ((1/D_B)/\mu_p(i))^\beta$ anneals from variance-reduction to unbiased estimation as β rises to 1.

3.2 The VLM as posterior, and semantic PER

Let $t^*(\tau)$ denote the unobserved outcome-causing timestep of trajectory τ , and $q_\phi(t^*|\tau)$ the VLM's approximate posterior. Define the per-transition semantic boost

$$w_{sem}(i; \tau) = \sum_{t^*} q_\phi(t^*|\tau) \cdot K_W(i; t^*) \quad (1)$$

where $K_W(i; t^*) = 1 + (w_{max} - 1) \cdot \mathbb{I}[|i - t^*| \leq W]$ is a window kernel of half-width W and bounded boost w_{max} . The semantic-PER sampling distribution is then the headline update

$$\mu_{sem}(i) \propto \mu_p(i) \cdot w_{sem}(i; \tau_i) \quad (2)$$

Two distinct factors drive priority: a *learning-progress* factor μ_p and a *causal-influence* factor w_{sem} . Standard PER uses only the first; uniform replay uses neither; semantic PER multiplies the two. When the VLM is non-committal the product degenerates to PER ($w_{sem} \approx 1$); when the VLM is confident, the failure window is up-weighted by w_{max} and TD-error further discriminates within it. By contrast VLM-RB [2] uses an *additive* mixture $\mu_{Sharony}(i) = \lambda \cdot \mu_{VLM}^P(i) + (1-\lambda) \cdot \mu_U(i)$ with a uniform—not TD-error—backbone, requiring a warm-up schedule $\lambda: 0 \rightarrow 0.5$ to avoid corrupting early training.

3.3 Off-policy correction taxonomy

Method	Proposal	Correction	Bias source
Retrace [10]	Behaviour policy μ_b	Clipped per-step IS c_s	IS-ratio truncation
V-trace [11]	Behaviour policy μ_b	Doubly-clipped IS ratios	Both truncations

Prioritised [1]	DQN	$\mu_p \propto \delta ^\alpha$	$w_{IS} = (D_B \cdot \mu_p)^{-\beta}$	Anneal of $\beta < 1$
Semantic (ours)	PER	$\mu_p \cdot w_{sem}$	$w_{IS,P}$ (under-corrected)	Miscalibration of q_ϕ
VLM-RB [2]		$\lambda \mu_{VLM}^P + (1-\lambda)\mu_U$	None published	Objective differs from PER's

Table 1. IS-corrected off-policy taxonomy. Semantic PER is a trajectory-level proposal-shaping scheme whose modifier q_ϕ is exogenous and foundation-model-supplied.

3.4 Exogenous credit-assignment oracle

The credit-assignment survey of Pignatelli et al. [23] groups methods by data source: endogenous TD-error (PER), endogenous return decomposition (RUDDER [24]), endogenous counterfactual baselines (CCA-PG [25]), endogenous counterfactual masking of LLM tokens (Khandoga et al. [26], CoSo [19]). A pre-trained VLM that answers ‘*which step caused the failure?*’ is, formally, an **exogenous foundation-model oracle**—a learned credit-assignment function pre-trained on a different distribution (web video, image-text pairs) and queried zero-shot. To our knowledge this is the first credit-assignment oracle entirely external to the RL training loop. Diagnostic predictions of this framing: semantic PER’s lift over PER should be monotone in (i) posterior concentration of q_ϕ and (ii) alignment of q_ϕ with the (privileged-state) Oracle posterior. A falsifier (adversarial-prompt experiment forcing q_ϕ to localise at random non-causal steps) is feasible at sweep budget and pre-registered in the supplementary plan.

4. Method

4.1 Counterfactual prompting analysis

We prototyped three prompt variants for the VLM at the failure timestep K: (A) *narrative*—free-form corrective description; (B) *action*—a 4D delta-action $[\Delta x, \Delta y, \Delta z, \text{grip}]$; (C) *achieved_goal*—a corrective 3D position. Aggregated over six failed Fetch episodes with Claude Opus 4.7 as both generator and independent judge, the narrative variant scored highest plausibility (0.79 mean, 4/4 plausible), the action variant scored highest specificity (0.65) but lowest confidence (0.43, with sign-flip errors on ~50% of y-axis components), and the achieved_goal variant was strongly task-dependent: mean 0.75 on FetchPickAndPlace (mid-air target) but 0.18 on FetchPush (table-surface target)—the VLM collapsed to *corrective_position* $\equiv$ *desired_goal*, asking the buffer to **teleport** the block. A head-to-head against a second model (Claude Sonnet 4.5, GPT-4o results in flight) confirms the teleport-collapse generalises across VLM families. Figure 2 reports the cross-model comparison.




Figure 2. Cross-model comparison of counterfactual prompt variants (C1v2). Sonnet 4.5 reproduces Opus’s teleport-collapse on Push and non-trivial mid-air corrective positions on PickAndPlace, confirming the failure mode is model-agnostic. GPT-4o data are in flight (insufficient_quota at run time); pipeline parameterised for swap-in.

4.2 VLM-Verified Counterfactual Hindsight

The teleport-collapse motivates a stronger acceptance gate. We change the prompted output to a corrective action sequence $a_{K:K+N}$ and **verify** it by snapshotting the MuJoCo state $s_K = (qpos, qvel, t, mocap, g)$ at the failure step, rolling the action sequence forward in a *separate* Gymnasium-Robotics env, and admitting the counterfactual only when the env’s own sparse reward fires (distance < 0.05 m on Fetch). A teleport instruction has no representation as a 4-vector action, so the failure mode is *structurally inexpressible*. On a 4/4 smoke test the verifier (i) rejects the teleport counterfactual with reason *no_corrective_action_provided*, (ii) evaluates VLM action counterfactuals end-to-end (final distance 0.16–0.32 m, rejected), and (iii) accepts a handcrafted close-goal corrective at 3.0 cm. The verifier costs 17.6 ms per call (50-step MuJoCo rollout, no rendering)—~0.4% training overhead on a 100k-step SAC run with 85% failed-episode rate, no extra VLM calls beyond the

existing counterfactual prompt. Round-trip snapshot tests pass exactly across FetchPush, FetchPickAndPlace, FetchSlide and FetchReach.

Theoretical novelty. This is model-based hindsight relabelling with a *perfect world model* (the simulator itself), differing from MuZero/Dreamer (learned model, modelling error) and from model-based HER (random rollouts) along three axes: the action source is a language prior, the search is N=1 single-rollout, and the acceptance criterion is the env's own sparse reward. The marriage of language prior with physics verification is to our knowledge novel: VLM-as-action-proposer (RT-2, OpenVLA) treats the VLM as policy; simulator-in-the-loop work uses sim for imagined-rollouts not verification; HER variants relabel without a verification gate. The three-way intersection is open territory.

5. Experiments

5.1 Setup

We train SAC [27] with HER and replay variants on the Gymnasium-Robotics Fetch suite: *FetchPush-v4*, *FetchPickAndPlace-v4*, *FetchSlide-v4*. Policy and critic are 3-layer MLPs on the proprioceptive vector state; RGB renders are computed only for the VLM evaluator. Buffer size 10^6 ; PER $\alpha = 0.6$, β anneals $0.4 \rightarrow 1.0$; semantic boost $w_{\max} = 10$, window $W = 5$. The VLM is Claude Opus 4.7 (with Sonnet 4.5 cross-check); a heuristic **Oracle** baseline reads privileged sim state (ee-object distance, object-goal distance, contact phases) to produce a ground-truth failure window, giving an upper-envelope curve the VLM result can be reported as a fraction of. We compare {Uniform, PER, PER+HER, Semantic PER (heuristic), Semantic PER (VLM), Bidirectional}.

5.2 Headline ablation (existing data)




Figure 3. Final success rate by method across FetchPush, FetchPickAndPlace and FetchSlide. Semantic PER (heuristic localiser, Oracle v2/v3) recovers a meaningful fraction of the PER+HER ceiling on contact-rich tasks. Mean $\pm$ SE across seeds; see Figure 4 for the aggregated cross-environment view.




Figure 4. Aggregated final success rates across the three Fetch tasks. Bidirectional (Section 5.5) tests whether failure and success signals are complementary.

5.3 Counterfactual prompts (C1, C1v2)

Per-variant judge scores ($n = 6$ failed episodes, six per cell): narrative 0.79 / 0.68 / 0.40 (plaus / progress / specificity), action 0.55 / 0.53 / 0.65, achieved_goal 0.46 / 0.97 / 0.49. The unified *all* prompt (narrative + action + position) **regularises away the teleport-collapse**: Push position deviates non-zero from desired_goal when paired with an action. The cross-model run (Claude Sonnet 4.5; GPT-4o in flight) confirms the failure mode is family-agnostic. We adopt the unified prompt with a 5 cm teleport-rejection radius and a $\text{sign}(\text{action}[3]) = \text{sign}(\text{desired} - \text{achieved})$ filter.

5.4 Cross-task transferability (N2)

We run the *same prompt template* (one Python format-string with *task_description* interpolated) on 4 failed episodes per env across {Push, PickAndPlace, Slide}. Headline result: 12/12 parse rate, 12/12 task-relevant failure annotations (independent Claude Opus 4.7 judge with env-specific rubric), 12/12 physically plausible. Tolerant keyframe agreement with the heuristic Oracle improves monotonically with failure-mode visual distinctiveness (50% Push $\rightarrow$ 75% PickAndPlace $\rightarrow$ 100% Slide); the Push number is a known artefact of the Oracle collapsing to $t = 0$ on synthetic stochastic rollouts. The contrast with any *learned* priority head (TD-error on a SAC checkpoint, contact-energy [5], or a fitted classifier) is structural: a learned head encodes the value-function geometry of one policy in one env and does not survive an env swap, whereas the VLM signal is task-agnostic by virtue of being *prompted* rather than learned.




Figure 5. Cross-task transferability of the VLM failure signal across FetchPush, FetchPickAndPlace and FetchSlide. Same prompt template (interpolated task description only), Claude Opus 4.7 as generator and judge. $n = 4$ episodes per env. Bars: mean $\pm$ SE.

5.5 In-flight: HER baselines and verification runs

Twenty-seven additional runs ($3 \text{ methods} \times 3 \text{ envs} \times 3 \text{ seeds}$) are in flight on Modal under tag *her_baselines*: {HER, HER+PER, HER+Semantic PER (heuristic)}. A bidirectional buffer variant (BidirectionalSemanticBuffer, Section 4.3 of the supplementary) boosts both failure-localised and success-localised transitions, isolating the incremental value of the failure-direction signal against Sharony's success-direction signal. Two 50k-step verification runs of the contact-aware Oracle v3 + bidirectional buffer (FetchPush, FetchSlide) are detached on Modal. Final numbers will populate Table 2 in the camera-ready.

6. Discussion, Limitations, and Future Work

Limitations. (i) The framing in Section 3 motivates rather than derives Semantic PER: the VLM posterior is not Bayesian-calibrated and we use the standard PER IS weight rather than the matching semantic IS weight, accepting a bias analogous to V-trace's truncation bias. (ii) The cross-task transferability result is from a 12-episode pilot; the Section 5.5 50-episode-per-env replication is in flight. (iii) GPT-4o head-to-head data are in flight (insufficient_quota at the cut-off); Sonnet 4.5 is the model-family stand-in. (iv) During development we discovered an unescaped-JSON bug in *src/vlm/localizer.py* (a *str.format* on a template containing literal JSON braces silently falls back to the exception handler); the cross-task runner works around it but the fix should land before the camera-ready.

Kill-or-commit experiments. Three pre-registered tests would falsify the IS-posterior framing: (a) adversarial prompt forcing q_ϕ to a random non-causal step; (b) modified-game ablation in the spirit of [2] (misleading sprites, abstract textures); (c) IS-correction ablation ($w_{IS,P}$ vs. $w_{IS,Sem}$) to quantify the under-correction. The cross-task scaling test ($N = 50/\text{env}$) and the addition of FetchReach + an OOD task (Adroit-Hand-Door) stress the transferability claim.

Future work. Soft-posterior prompting (VLM emits a distribution over t^*) would let us use the exact expectation in Eq. (1) and monitor effective sample size; freshness-aware decay [12] would address the policy-drift coupling of q_ϕ to the buffer; the combined *Sharony* $\times$ *Semantic-PER* ablation tests whether success- and failure-direction signals are orthogonal; the verified-CF buffer subclass closes the integration with C2's stub for an end-to-end VLM-verified hindsight relabelling pipeline.

References

- [1] T. Schaul, J. Quan, I. Antonoglou, and D. Silver. Prioritized Experience Replay. ICLR, 2016.
- [2] E. Sharony, T. Jurgenson, O. Krupnik, D. Di Castro, and S. Mannor. VLM-Guided Experience Replay. arXiv:2602.01915, 2026.
- [3] M. Andrychowicz et al. Hindsight Experience Replay. NeurIPS, 2017.
- [4] F. Özgür, R. Zurbrugg, and S. Kumar. Next-Future: Sample-Efficient Policy Learning for Robotic-Arm Tasks. arXiv:2504.11247, 2025.
- [5] E. Sayar, Z. Bing, C. D'Eramo, O. S. Oguz, and A. Knoll. Contact Energy Based Hindsight Experience Prioritization. arXiv:2312.02677, 2023.
- [6] C. Chuck, F. Feng, C. Qi, C. Shi, S. Agarwal, A. Zhang, and S. Niekum. Null Counterfactual Factor Interactions for Goal-Conditioned RL (HInt/NCII). ICLR, 2025.
- [7] X. Lei et al. GCHR: Goal-Conditioned Hindsight Regularization. arXiv:2508.06108, 2025.
- [8] Y. Zhang and K. Mason. D-SPEAR: Dual-Stream Prioritized Experience Adaptive Replay. IEEE ICCRE, 2026.
- [9] P. Zhou et al. Act2Goal: From World Model to General Goal-conditioned Policy. arXiv:2512.23541, 2025.
- [10] R. Munos, T. Stepleton, A. Harutyunyan, and M. G. Bellemare. Safe and Efficient Off-Policy Reinforcement Learning (Retrace). NeurIPS, 2016.
- [11] L. Espeholt et al. IMPALA / V-trace: Scalable Distributed Deep-RL with Importance-Weighted Actor-Learner Architectures. ICML, 2018.
- [12] W. Ma et al. Freshness-Aware Prioritized Experience Replay for LLM/VLM Reinforcement Learning. arXiv:2604.16918, 2026.
- [13] J. Rocamonde, V. Montesinos, E. Nava, E. Perez, and D. Lindner. Vision-Language Models are Zero-Shot Reward Models for RL. ICLR, 2024.
- [14] O. Y. Lee, A. Xie, K. Fang, K. Pertsch, and C. Finn. Affordance-Guided RL via Visual Prompting (KAGI). IROS, 2025.
- [15] D. Venuto et al. Code as Reward: Empowering RL with VLMs. ICML, 2024.
- [16] Y. Wu et al. Large Reward Models: Generalizable Online Robot Reward Generation with VLMs. arXiv:2603.16065, 2026.
- [17] W. Chen, O. Mees, A. Kumar, and S. Levine. Vision-Language Models Provide Promptable Representations for RL. arXiv:2402.02651, 2024.
- [18] C. Glossop, W. Chen, A. Bhorkar, D. Shah, and S. Levine. CAST: Counterfactual Labels Improve Instruction Following in VLA Models. arXiv:2508.13446, 2025.
- [19] L. Feng et al. Towards Efficient Online Tuning of VLM Agents via Counterfactual Soft RL (CoSo). ICML, 2025.
- [20] J. Duan et al. AHA: A Vision-Language-Model for Detecting and Reasoning Over Failures in Robotic Manipulation. ICLR, 2025.
- [21] L. Ding. AgentHER: Hindsight Experience Replay for LLM Agent Trajectory Relabeling. arXiv:2603.21357, 2026.
- [22] M. Y. Hu, B. Van Durme, J. Andreas, and H. Jhamtani. Sample-Efficient Online Learning in LM Agents via Hindsight Trajectory Rewriting (ECHO). arXiv:2510.10304, 2026.
- [23] E. Pignatelli et al. A Survey of Temporal Credit Assignment in Deep Reinforcement Learning. arXiv:2312.01072, 2023.
- [24] J. A. Arjona-Medina et al. RUDDER: Return Decomposition for Delayed Rewards. NeurIPS, 2019.
- [25] T. Mesnard et al. Counterfactual Credit Assignment in Model-Free RL (CCA-PG). ICML, 2021.
- [26] M. Khandoga, R. Yuan, and V. K. Sankarapu. Beyond Uniform Credit: Causal Credit Assignment for Policy Optimization. arXiv:2602.09331, 2026.
- [27] T. Haarnoja, A. Zhou, P. Abbeel, and S. Levine. Soft Actor-Critic: Off-Policy Maximum Entropy Deep RL with a Stochastic Actor. ICML, 2018.